AIHackr — Execution Playbook

For: aihackr.com
Compiled: April 2026
Covers: Repositioning · Recurring-Use Hook · Content & Leaderboard Engine

> This document is the complete execution plan for the three highest-leverage moves to find PMF. It is written so a developer and a marketer can start Monday morning without further clarification.

Table of Contents

1. Part 1: Repositioning Around the AI-Provider Wedge
2. Part 2: Recurring-Use Hook — Watchlists, Alerts & Digest
3. Part 3: Content & Leaderboard Engine

AIHackr Repositioning Guide

From "SaaS tech stack detection" to "the only tool that tells you which AI brain a SaaS product thinks with"

Current State Diagnosis

The homepage currently reads: *"What's this SaaS built with? Especially the AI stack."* The badge says *"The 'BuiltWith' for AI SaaS."* Both are self-inflicted wounds. Calling yourself the BuiltWith for anything invites the obvious question: why not just use BuiltWith? Positioning as a general tech-stack detector puts you in a race you cannot win — BuiltWith has 20+ years of domain data and Wappalyzer has a browser extension installed on millions of devices. You will lose on framework detection. You will lose on breadth. You will never win on volume.

The fix is not to add more tech categories. The fix is to own the one category where you have a genuine, defensible moat: **AI provider intelligence** — knowing whether a product runs on GPT-4, Claude, Gemini, or a fine-tuned model, with a confidence score and an evidence trail that holds up in a board meeting or a due diligence memo.

No one else does this. That is your entire positioning in one sentence.

1. Messaging Hierarchy

The Single Core Message

AIHackr tells you which AI provider is powering any SaaS product — with confidence levels and evidence — so you stop guessing and start knowing.

Primary Value Prop (2–3 sentences)

Every serious AI startup claims to be "AI-native" and "model-agnostic." AIHackr cuts through the marketing to show you exactly which LLM provider a product is running — OpenAI, Anthropic, Gemini, Azure OpenAI — and how confident that signal is. For VCs doing diligence, PMs tracking competitor migrations, and founders benchmarking their own infrastructure decisions, that's the intelligence layer that hasn't existed until now.

3 Supporting Proof Points

1. **No one else has this signal.** BuiltWith tells you the frontend framework. Wappalyzer tells you the analytics vendor. Neither tells you whether a product is running GPT-4 Turbo, Claude 3.5 Sonnet, or a fine-tuned open-source model — AIHackr does.
2. **Evidence trails, not gut checks.** Every detection comes with the specific technical signals that produced the confidence score — response headers, API call patterns, error signatures, and behavioral fingerprints. You can show your work in a memo or a board deck.
3. **The AI layer is the one that matters now.** In 2024–2025, the most strategically important question about any AI product is not what framework it runs on — it is which model it trusts with its core reasoning. That's the due diligence question. That's the competitive intelligence question. That's the question AIHackr answers.

What to Cut from Current Messaging

- **The "BuiltWith" comparison in the badge.** It frames you as a follower, not a category creator. Cut it immediately.
- **"Reverse-Engineer Any SaaS Stack"** as the core frame. "Any" implies breadth. Breadth is where you lose. Specificity is where you win.
- **The laundry list subheadline.** "framework, hosting, payments, auth, analytics, and AI provider signals" buries the lede. AI provider signals are not an item in a list — they are the entire product.
- **"Shareable Stack Card"** as the primary artifact name. It undersells. Call it what it is: an AI Intelligence Report or an AI Provider Report.
- **"Progressive scan: instant results from headers & HTML, then deeper browser analysis runs automatically in the background."** This is a technical explanation of your delivery mechanism. No one cares how the scan works. They care what they learn. Move this to a FAQ or technical docs.

2. Homepage Copy Rewrites

Hero Headline — 3 A/B Test Variants

Variant A (Clarity):

> Which AI Is This SaaS Actually Running?

Variant B (Contrast / Intrigue):

> They Say "AI-Native." We'll Tell You Which One.

Variant C (Direct to ICP):

> Find Out If Your Competitor Just Switched from OpenAI to Claude.

Hero Subheadline

> Paste a URL. In seconds, AIHackr identifies the AI provider powering that product — OpenAI, Anthropic, Gemini, Azure OpenAI — with a confidence score and the specific signals that produced it. It's the intelligence layer that didn't exist before due diligence meetings, competitive reviews, or investor calls. Now it does.

Primary CTA Button Text

> Scan a Competitor →

(Alternatives to A/B test: "Detect Their AI Stack" / "Find Out Now")

Secondary CTA

> See a Sample AI Intelligence Report

(Replaces the current generic "See a sample report" — name the artifact so the CTA sets expectations)

Social Proof Line (Testimonials Section Header)

> What VCs, PMs, and founders are finding out for the first time.

(Individual testimonial format: Lead with the insight, not the endorsement. Example: "[Company X] had "GPT-4" in all their marketing. AIHackr showed Azure OpenAI with 91% confidence. Changed how we valued the moat." — Partner, [Firm]

Comparison Table Header Copy

Reframed table header:

> The intelligence BuiltWith and Wappalyzer don't have — and can't get.

Table intro copy:

> They track your CSS framework. We track your intelligence layer.

What you want to know	BuiltWith	Wappalyzer	AIHackr
Frontend framework	✓	✓	✓
Hosting provider	✓	✓	✓
Analytics tools	✓	✓	✓
Which AI/LLM provider	✗	✗	✓
Confidence score + evidence	✗	✗	✓
OpenAI vs Claude vs Gemini vs Azure	✗	✗	✓
Fine-tune vs. base model signals	✗	✗	✓

Table footer:

> BuiltWith and Wappalyzer are excellent at what they do. This is the thing they don't do.

Feature Section Headers (3 Features, AI-Detection-Forward)

Feature 1:

> AI Provider Fingerprinting
> We detect OpenAI, Anthropic, Gemini, and Azure OpenAI from behavioral patterns, API signatures, and response characteristics — the signals that are hard to hide even when companies don't want you to know.

Feature 2:

> Confidence Scores With Evidence Trails
> Every detection comes with a confidence level (High / Medium / Low) and the specific signals that produced it. Not a guess. A documented finding you can put in a memo, a deck, or a diligence report.

Feature 3:

> Migration Monitoring — Know When They Switch
> Set up alerts for any domain. When a competitor migrates from GPT-4 to Claude, or moves from OpenAI's API to Azure OpenAI for compliance reasons, you'll know before their next press release.

3. Meta & SEO Copy

Page Title Tag *(55 chars)*

```
AIHackr — Detect Which AI Provider Any SaaS Uses
```

Meta Description *(152 chars)*

```
AIHackr detects which AI/LLM provider powers any SaaS — OpenAI, Claude, Gemini, or A
```

(Note: trim to 155 chars if rendering long — the above is 160 chars exactly. Suggested trim: drop "and evidence")

Trimmed version (143 chars):

```
AIHackr detects which AI/LLM provider powers any SaaS — OpenAI, Claude, Gemini, or A
```

OG Title

```
AIHackr: Which AI Is That SaaS Actually Running?
```

OG Description

```
Paste a URL. Find out if a product runs on OpenAI, Anthropic, Gemini, or Azure Open
```

Twitter Card Copy

Title:

```
AIHackr — AI Provider Detection for Any SaaS
```

Description:

```
Stop guessing which LLM a competitor is running. AIHackr detects OpenAI, Claude, Ge
```

4. Social & Profile Copy

Twitter/X Bio *(158 chars)*

```
Detect which AI provider (OpenAI, Claude, Gemini, Azure) is powering any SaaS — witl
```

LinkedIn Company Description *(298 chars)*

```
AIHackr is the only tool that tells you which AI/LLM provider is powering any SaaS pr
```

ProductHunt Tagline

```
Detect which AI provider any SaaS is actually running
```

(Alternate: "BuiltWith can see the framework. We can see the brain.")

5. The Narrative: Talking Points

These are for interviews, podcasts, Twitter threads, and investor conversations. Opinionated, quotable, designed to create contrast.

Talking Point 1 — The Category Claim

> "Every AI startup says they're 'AI-native' and 'model-agnostic.' That's marketing. The real question — the one that matters for due diligence, for competitive intelligence, for understanding whether a company has real AI infrastructure or just an OpenAI API key wrapped in a pretty UI — is: which model is actually doing the reasoning? BuiltWith can tell you they use React. We can tell you they think in Claude. That's a different kind of intelligence."

Talking Point 2 — The VC / Diligence Frame

> "I've watched VCs spend weeks on technical diligence trying to figure out if an AI company has proprietary model infrastructure or is just a thin wrapper around GPT-4. That question now has a three-second answer. Paste the URL. We tell you the provider, the confidence level, and the evidence. That's not a party trick — that changes how you evaluate defensibility."

Talking Point 3 — The Migration Signal

> "The most underrated competitive intelligence signal in AI right now isn't product announcements or pricing pages — it's provider migrations. When a company quietly moves from OpenAI to Anthropic, or from the OpenAI API to Azure OpenAI, that tells you something. Maybe they're optimizing for cost. Maybe they're meeting an enterprise compliance requirement. Maybe their GPT-4 results weren't good enough for their use case. We track those migrations. Most companies don't even know they're broadcasting them."

Talking Point 4 — The Contrarian Take on 'AI-Powered'

> "The phrase 'AI-powered' is the new 'cloud-based.' It means everything and nothing. There are companies running Claude 3.5 Sonnet on custom fine-tuned datasets and companies with a single if-statement calling the GPT-3.5 API — and both of them have 'AI' in their pitch deck. The difference in capability, cost structure, and defensibility is enormous. AIHackr is for people who want to know which side of that line a company is actually on."

Talking Point 5 — The Market Timing Angle

> "Two years ago, nobody cared which database a SaaS ran on. Today, which AI provider a company uses is a first-order strategic question. It affects cost structure, reliability, model capability, compliance posture, and vendor lock-in risk. The infrastructure layer that matters has shifted from cloud providers to AI providers — and there's been no way to see it from the outside. Until now."

6. What To Kill

These are the five specific elements on the current AIHackr site and in its messaging that directly undermine the AI-provider positioning. Remove or deprioritize all of them.

Kill 1: The "BuiltWith" badge

Current: The hero section badge reads *"The 'BuiltWith' for AI SaaS"*

Why it has to go: It's the single most damaging line on the site. It invites users to compare you to BuiltWith on BuiltWith's terms — breadth of tech detection — and you will lose that comparison every time. It also positions you as a derivative ("the X for Y") rather than a category creator. Your product does something BuiltWith categorically cannot do. Lead with that. Don't define yourself through them.

Kill 2: The comma-separated feature list in the subheadline

Current: "framework, hosting, payments, auth, analytics, and AI provider signals when publicly detectable"

Why it has to go: This is a capabilities dump, not a value proposition. The phrase "when publicly detectable" is also a hedging qualifier that weakens the whole sentence. Listing AI detection as item six in a seven-item list makes it look like an afterthought on a general-purpose tool. It is not an afterthought. It is the product.

Kill 3: The "Progressive scan" technical explanation below the CTA

Current: "Progressive scan: instant results from headers & HTML, then deeper browser analysis runs automatically in the background."

Why it has to go: This is implementation documentation placed at the highest-attention moment of the page — immediately below the primary CTA. No one who is about to paste a URL into a scan box needs to know how the scan pipeline works. This copy answers a question no prospect was asking. Replace it with a trust signal: number of companies scanned, detection accuracy rate, or a one-line social proof quote.

Kill 4: "Reverse-Engineer Any SaaS Stack" as a positioning frame

Why it has to go: "Any" is the problem word. It positions AIHackr as trying to do everything BuiltWith and Wappalyzer already do, but for SaaS. You cannot win on "any." You win on "the one thing no one else can tell you." Also, "reverse-engineer" has a slightly adversarial, hacking connotation that might spook enterprise buyers or compliance teams. Replace it with clear, confident language about AI provider intelligence.

Kill 5: Generic tech stack breadth as a selling point

Why it has to go: If any part of your site, landing pages, or marketing is leading with the full breadth of tech detection (hosting providers, CSS frameworks, payment processors, auth tools), you are diluting the brand and confusing the ICP. A VC doing AI startup diligence does not care that you can detect Stripe. A PM tracking a competitor's AI infrastructure does not care about Cloudflare detection. Every byte of homepage real estate spent on generic tech-stack breadth is a byte not spent on the one capability that makes someone choose AIHackr over any other tool. Deprioritize breadth entirely. It can live in a features page or documentation, not the homepage.

Positioning Summary Card

(One-pager version for quick reference)

Element	Current	New
Category	SaaS tech stack detector	AI provider intelligence
Core claim	Reverse-engineer any SaaS stack	Find out which AI brain it thinks with
Comparison frame	"The BuiltWith for AI SaaS"	"The intelligence BuiltWith doesn't have"
Primary ICP	Developers / curious builders	VCs, competitive PMs, AI-native founders
Key differentiator	General tech detection	AI/LLM provider detection + confidence + evidence
Tone	Tool / utility	Intelligence layer / strategic weapon

Guide prepared April 2025. Competitor positioning verified against live sites for BuiltWith (builtwith.com) and Wappalyzer (wappalyzer.com).

AIHackr — Watchlist & Alerts Product Spec

Document type: Feature PRD
Product: AIHackr.com
Feature set: AI Stack Watchlists + Change Alerts + Weekly Digest
Status: Draft v1.0
Purpose: Recurring-use hook to drive retention and paid conversion

Background

AIHackr detects which AI/LLM providers (OpenAI, Anthropic, Gemini, Azure OpenAI, Cohere, Mistral, etc.) a SaaS product is using, with confidence levels derived from network signals, SDK fingerprints, and API call patterns. The core retention problem: a tech stack lookup is inherently a one-time action. Without a reason to return, a user who runs a scan once has no pull to come back.

The solution is a recurring-use loop built on three tightly coupled features:

1. **AI Stack Watchlists** — users save domains and AIHackr monitors them automatically
2. **Change Alerts** — push notifications (email + Slack) when a monitored domain's AI stack changes
3. **Weekly AI Stack Digest** — a curated, data-driven email surfacing what changed across a user's watchlist and what's trending market-wide

Together, these transform AIHackr from a lookup tool into an intelligence subscription — the mechanism that justifies a monthly price.

1. Feature Overview

AI Stack Watchlists

An AI Stack Watchlist lets users save a set of domains they care about — competitors, prospects, portfolio companies, industry peers — and have AIHackr continuously monitor their AI provider footprint. Instead of manually re-running scans, users maintain a live dashboard showing the current AI stack for every watched domain, when it was last scanned, what providers are active, and whether anything has changed since the last check. The watchlist is the product's persistent value surface: a user who has saved 20 domains and sees them updating automatically has a reason to return every week and has already built switching costs.

Change Alerts (Email + Slack)

Change Alerts fire whenever a monitored domain's AI stack shifts in a meaningful way — a new provider detected, an existing provider removed, or a confidence level crossing a defined threshold. Alerts are delivered via email and optionally via a Slack webhook, formatted as a concise signal with the specific change, the affected domain, the before/after provider state,

and a direct link to the full scan result. Users can configure alert sensitivity (e.g., only fire on high-confidence changes, or only for specific providers). The goal is to make AIHackr the fastest way to learn that a competitor just switched from OpenAI to Anthropic — turning a passive scan tool into an active competitive intelligence feed.

Weekly AI Stack Digest

The Weekly AI Stack Digest is a scheduled email delivered every Monday morning that summarizes the prior week's activity across a user's watchlist. It shows which domains changed, what their stacks look like now, and surfaces market-level trends aggregated anonymously across all AIHackr users' watchlists — e.g., "Anthropic adoption up 14% among B2B SaaS tools this week" or "3 of your 12 watched competitors added Azure OpenAI." The digest is designed to be a 60-second read that delivers enough signal to forward to a team, link in a Slack channel, or file as competitive intel — creating organic sharing and word-of-mouth acquisition.

2. Watchlist Feature Spec

2.1 Data Model

`watchlist_entry` table

Field	Type	Description
`id`	UUID	Primary key
`user_id`	UUID FK	Owner of this watchlist entry
`domain`	VARCHAR(255)	Normalized domain (e.g., `linear.app`) — no protocol, no trailing slash
`label`	VARCHAR(100)	Optional user-defined label (e.g., "Competitor", "Prospect")
`added_at`	TIMESTAMP	When user added this domain
`last_scanned_at`	TIMESTAMP	When AIHackr last ran a scan on this domain
`next_scan_at`	TIMESTAMP	Scheduled next scan time
`scan_status`	ENUM	`pending`, `scanning`, `complete`, `error`, `unreachable`
`is_active`	BOOLEAN	User can pause monitoring without deleting
`alert_enabled`	BOOLEAN	Whether change alerts are on for this entry
`alert_threshold`	ENUM	`any_change`, `high_confidence_only`, `provider_added_removed`
`notes`	TEXT	Optional user notes

`scan_result` table (one record per scan run per domain)

Field	Type	Description
`id`	UUID	Primary key
`watchlist_entry_id`	UUID FK	Parent watchlist entry
`domain`	VARCHAR(255)	Domain scanned
`scanned_at`	TIMESTAMP	Scan execution timestamp
`scan_duration_ms`	INT	Time taken
`providers_detected`	JSONB	Array of provider detection objects (see below)
`raw_signals`	JSONB	Raw evidence dict (SDK names, headers, endpoints seen)
`site_reachable`	BOOLEAN	Whether domain responded
`http_status`	INT	HTTP status code of main domain response
`error_message`	TEXT	Error detail if scan failed
`scan_version`	VARCHAR(20)	Detection engine version (for reproducibility)

`provider_detection` object (embedded in `scan_result.providers_detected`)

```
{
  "provider": "openai",
  "display_name": "OpenAI",
  "confidence": "high",          // "high" | "medium" | "low"
  "confidence_score": 0.91,      // float 0.0-1.0
  "signals": ["sdk_detected", "api_endpoint_pattern", "response_header"],
  "model_hints": ["gpt-4o", "gpt-4-turbo"],  // optional, if detectable
  "first_detected_at": "2025-01-10T08:00:00Z",
  "last_seen_at": "2025-01-17T09:15:00Z"
}
```

`stack_change` table (one record per detected change)

Field	Type	Description
`id`	UUID	Primary key
`watchlist_entry_id`	UUID FK	Which watchlist entry changed
`scan_result_id`	UUID FK	The scan that detected the change
`prior_scan_result_id`	UUID FK	The previous scan being compared against
`change_type`	ENUM	`provider_added`, `provider_removed`, `confidence_upgraded`, `confidence_downgraded`, `model_hint_changed`
`provider`	VARCHAR(50)	Affected provider (e.g., `anthropic`)

Field	Type	Description
`prior_confidence`	ENUM	`high`, `medium`, `low`, `none`
`new_confidence`	ENUM	`high`, `medium`, `low`, `none`
`change_severity`	ENUM	`major`, `minor` — used for alert throttling
`alerted_at`	TIMESTAMP	When alert was sent (null if not yet sent or suppressed)
`alert_suppressed`	BOOLEAN	True if change was below user's alert threshold
`detected_at`	TIMESTAMP	When change was detected

2.2 UX Flow

Adding a Domain

1. From any scan result page, the user clicks "+ **Add to Watchlist**" (present in the top-right of every result card).
   — Alternatively, from the Watchlist dashboard, click "**Add Domain**" → type or paste a domain.
2. A modal appears with three fields:
   - **Domain** (pre-filled if coming from a scan result)
   - **Label** (optional, free text — e.g., "Competitor", "Prospect Q1")
   - **Alert threshold** (radio: "Any change detected" / "High-confidence changes only" / "Only provider additions or removals") — default: "Any change detected"
3. User clicks "**Start Monitoring**."
4. An immediate background scan is triggered. The entry appears in the watchlist with status `scanning…` and animates to `complete` within 15–60 seconds.
5. If the user is on the Free plan and has reached the watchlist limit (5 domains), the modal surfaces a plan upgrade prompt instead of saving.

Watchlist Dashboard Layout

The Watchlist page (`/watchlist`) is the product's home base for logged-in users. Layout:

Top bar:

- "My Watchlist" heading with domain count ("12 domains")
- "**Add Domain**" button (primary CTA)
- Filter chips: All | Changed This Week | High Confidence | Errors
- Sort dropdown: Last Changed | Alphabetical | Last Scanned

Main table (one row per domain):

Column	Description
Domain	Favicon + domain name + optional user label (in muted text below)
AI Providers	Provider logos with confidence color-coding (green = high, yellow = medium, gray = low). Up to 3 shown inline, "+ N more" overflow.
Confidence	Highest confidence level across all detected providers
Last Scanned	Relative timestamp ("2 hours ago") with tooltip for absolute time
Last Changed	Date of most recent stack change, or "No changes detected"
Change Summary	Badge showing what changed: e.g., "+ Anthropic" (green) or "- OpenAI" (red) or "Confidence drop: GPT-4" (orange)
Status	Icon: green checkmark (healthy), yellow spinner (scanning), red X (error/unreachable), gray pause (monitoring paused)
Actions	"View Scan", "Edit", "Pause/Resume", "Remove"

Row expansion (click any row): Expands inline to show:

- Full provider detection list with all signals
- Change history timeline (last 10 changes, sparkline)
- Link to full scan report

Empty state: If no domains in watchlist, show onboarding prompt: "Add your first competitor or prospect to start monitoring their AI stack."

2.3 Change Detection Logic

What Constitutes a "Change"

A change is any delta between the most recent scan result and the prior scan result for the same domain. The following events are classified as changes:

Event	Change Type	Severity
A provider is detected that was not in the prior scan	`provider_added`	Major
A provider present in the prior scan is no longer detected	`provider_removed`	Major
A provider's confidence moves from `low` or `medium` → `high`	`confidence_upgraded`	Minor
A provider's confidence moves from `high` → `medium` or `low`	`confidence_downgraded`	Minor
A model hint appears or changes (e.g., GPT-4o detected where GPT-3.5 was before)	`model_hint_changed`	Minor

Changes that are **NOT** recorded (noise suppression):

- Confidence score fluctuations of less than 0.10 within the same confidence band (e.g., 0.82 → 0.79 — both "high")

- Signal count fluctuations with no category-level change (still same providers at same confidence)
- Scan result differences attributable to a known engine version upgrade (flagged by `scan_version` mismatch — log but suppress alert; re-baseline)

Re-Scan Cadence

Plan	Re-scan interval
Free	Every 7 days
Pro	Every 24 hours
(Future) Enterprise	Every 6 hours

Re-scans run as background jobs via a scan queue. Priority: Pro domains are head-of-queue. A short random jitter (±2 hours) is added to avoid thundering herd on domain hosts. Manual "Scan Now" is available to Pro users (max 3× per domain per day).

Diff Algorithm

```
function computeDiff(prior_scan, new_scan):
  prior_providers = {p.provider: p for p in prior_scan.providers_detected}
  new_providers   = {p.provider: p for p in new_scan.providers_detected}

  changes = []

  // Provider added
  for provider in new_providers - prior_providers:
    if new_providers[provider].confidence_score >= 0.60:  // ignore very-low-confid
      changes.append(ChangeEvent(type=PROVIDER_ADDED, provider=provider, ...))

  // Provider removed
  for provider in prior_providers - new_providers:
    if prior_providers[provider].confidence_score >= 0.60:  // don't alert on remov
      changes.append(ChangeEvent(type=PROVIDER_REMOVED, provider=provider, ...))

  // Confidence band shift for shared providers
  for provider in prior_providers ∩ new_providers:
    prior_band = confidenceBand(prior_providers[provider].confidence_score)
    new_band   = confidenceBand(new_providers[provider].confidence_score)
    if prior_band != new_band:
      change_type = CONFIDENCE_UPGRADED if new_band > prior_band else CONFIDENCE_DOW
      changes.append(ChangeEvent(type=change_type, provider=provider, ...))

  // Model hint changes (minor)
  for provider in prior_providers ∩ new_providers:
    if prior_providers[provider].model_hints != new_providers[provider].model_hin
      changes.append(ChangeEvent(type=MODEL_HINT_CHANGED, ...))

  return changes

function confidenceBand(score):
  if score >= 0.80: return HIGH    // 3
```

```
    if score >= 0.50: return MEDIUM  // 2
    return LOW                       // 1
```

Changes are written to the `stack_change` table. Alert delivery is a separate async step that reads from this table and filters by user alert threshold settings.

2.4 Edge Cases

Provider Not Detectable (No Signal)

- **Behavior:** Store `providers_detected = []` in the scan result. Mark as "No AI providers detected" in UI with a muted "?" icon.
- **Change logic:** If prior scan had providers and new scan detects none, this is treated as `provider_removed` for each previously detected provider. This is a **Major** change and alerts fire.
- **Display:** Watchlist row shows "No providers detected" with last-seen-provider in parentheses: "No providers detected (previously: OpenAI)."

Site Goes Down During Scan

- **Behavior:** If HTTP status is 5xx, connection timeout, or DNS failure, mark `scan_status = error`, `site_reachable = false`. Do **not** write any `stack_change` records. Do **not** treat unreachability as a provider removal — this is a transient infrastructure event, not a stack change.
- **Retry:** Retry scan after 4 hours (up to 3 retries). If all retries fail, mark `scan_status = unreachable` and surface a warning badge in the watchlist row.
- **Alert:** Do not send change alerts. If site is unreachable for >3 consecutive scans (>3 days on Free, >3 days on Pro), send a one-time "Site unreachable" notification.

False Positive (Confidence Drop Without Actual Provider Change)

- **Scenario:** A site's detection signals are temporarily obscured (CDN A/B test, maintenance page, Cloudflare challenge) causing confidence to drop from High → Low for a provider that is still in use.
- **Mitigation 1:** Confidence downgrade alerts are classified as **Minor** severity. Users on "Provider added/removed only" threshold will not receive an alert.
- **Mitigation 2:** A single-scan confidence drop is not enough to record a `PROVIDER_REMOVED` event. Provider removal requires 2 consecutive scans with no detection (configurable, default: 2). This is enforced in the diff algorithm via a "pending removal" state in the watchlist entry.
- **Mitigation 3:** In the alert email and Slack message, confidence downgrade events are flagged with a note: "This may be a temporary detection change. We recommend verifying before acting."

- **User control:** Users can mark a specific change as "Dismissed / False Positive" in the UI, which suppresses future alerts for that specific provider/confidence-band combination for 30 days.

3. Change Alert Spec

3.1 Email Alert Template

Subject line formula:
`[AIHackr] {Domain} just {change_verb} {Provider}` ← for major changes
Examples:

- `[AIHackr] linear.app just added Anthropic`
- `[AIHackr] notion.so just dropped OpenAI`
- `[AIHackr] figma.com — AI stack change detected`

For minor changes (confidence shifts):
`[AIHackr] {Domain} — confidence update on {Provider}`

Email body:

```
From: AIHackr Alerts <alerts@aihackr.com>
Subject: [AIHackr] linear.app just added Anthropic

────────────────────────────────────
          AIHackr Stack Alert
────────────────────────────────────

AI stack change detected on linear.app

What changed
────────────
  ✅ NEW   Anthropic (Claude)        High confidence
  ──────────────────────────────────
  ✓ Still  OpenAI (GPT-4o)           High confidence
  ✓ Still  Azure OpenAI              Medium confidence

Detected: Monday, January 20, 2025 at 09:14 AM UTC
Prior scan: January 13, 2025

Evidence signals:
  • Anthropic SDK detected in frontend bundle
  • API endpoint pattern matching Anthropic messages API
  • Response latency profile consistent with Claude 3.5

────────────────────────────────────

[ View Full Scan Report → ]    [ View Watchlist → ]

────────────────────────────────────
Why this matters: Detecting a new LLM provider often
```

```
signals a new product feature, a model swap, or an
infrastructure change — usually within 2–4 weeks of
the deployment going live.

──────────────────────────────────────

You're receiving this because you added linear.app to
your AIHackr watchlist with alerts on.

[ Manage alert settings ]  |  [ Unsubscribe from alerts ]
AIHackr · aihackr.com
```

Key design principles for the email:

- Lead with the change, not the brand. The subject line is the signal.
- Show before/after state so the user doesn't need to click to understand what happened.
- Include the evidence signals — this is what separates a signal from spam. Users should be able to judge whether to act without visiting the site.
- The "Why this matters" blurb is static but contextual — it makes the email feel editorial, not automated.
- Single primary CTA ("View Full Scan Report") — secondary CTA is watchlist. No more than two CTAs.

3.2 Slack Notification Format

Delivered via user-configured Slack webhook (Bot or incoming webhook, no OAuth required for MVP).

```
⬛ *AI Stack Change Detected*

*linear.app* just added *Anthropic (Claude)*

*Before* → *After*
──────────────────────────────
✅ `ADDED`   Anthropic — ⬛ High confidence
✓  `STABLE`  OpenAI (GPT-4o) — ⬛ High confidence
✓  `STABLE`  Azure OpenAI — ⬛ Medium confidence

⬛ *Detected:* Jan 20, 2025 · 09:14 AM UTC
⬛ <https://aihackr.com/scan/abc123|View Full Scan Report>
⬛ <https://aihackr.com/watchlist|Open Watchlist>

> *Evidence:* Anthropic SDK in bundle · API endpoint match · Latency profile consist
```

For **minor changes** (confidence shifts), use a subdued format:

```
⚠ *Confidence update — linear.app*

OpenAI confidence shifted: ⬛ High → ⬛ Medium
This may be a temporary detection fluctuation.
```

```
 <https://aihackr.com/scan/abc123|View Scan>
```

Slack webhook payload shape:

```
{
  "text": "AI Stack Change: linear.app added Anthropic",
  "blocks": [
    { "type": "header", "text": { "type": "plain_text", "text": "  AI Stack Change [
    { "type": "section", "text": { "type": "mrkdwn", "text": "*linear.app* just add
    { "type": "section", "fields": [
      { "type": "mrkdwn", "text": "*Provider:* Anthropic (Claude)" },
      { "type": "mrkdwn", "text": "*Confidence:*  High" },
      { "type": "mrkdwn", "text": "*Detected:* Jan 20, 2025" },
      { "type": "mrkdwn", "text": "*Change type:* Provider added" }
    ]},
    { "type": "actions", "elements": [
      { "type": "button", "text": { "type": "plain_text", "text": "View Full Report"
      { "type": "button", "text": { "type": "plain_text", "text": "Open Watchlist" }
    ]}
  ]
}
```

3.3 Alert Settings

Users configure alerts at two levels: **global** (account-level defaults) and **per-domain** (overrides per watchlist entry).

Global alert settings (`/settings/alerts`):

Setting	Options	Default
Email alerts	On / Off	On
Slack alerts	On / Off (+ webhook URL)	Off
Alert frequency cap	Immediate / Max 1/day / Max 1/week per domain	Immediate
Change threshold	Any change / High-confidence only / Provider added or removed only	Any change
Minimum confidence to alert	Low / Medium / High	Medium
Quiet hours	No quiet hours / Custom time window (timezone-aware)	No quiet hours
Alert digest mode	Send individual alerts / Bundle into daily digest	Individual

Per-domain override (on the watchlist entry edit modal):

Setting	Options
Override global alert threshold	Use global / Any change / Provider added or removed only / No alerts

Setting	Options
Provider-specific filter	Alert only if change involves: [multiselect: OpenAI, Anthropic, Gemini, Azure, All]
Pause alerts	Pause for: 1 day / 1 week / Until I turn back on

Implementation note: Alert delivery is gated by a deduplication key: `(watchlist_entry_id, provider, change_type, date)`. The same change cannot fire more than once in a 24-hour window regardless of settings, preventing double-sends from retry logic.

4. Weekly Digest Email Spec

4.1 Subject Line Formula

`Your AI Stack Digest: {N} change{s} across {M} watched {domain/domains} this week`

Examples:

- `Your AI Stack Digest: 3 changes across 12 watched domains this week`
- `Your AI Stack Digest: linear.app added Anthropic — and 2 more changes`
- `Your AI Stack Digest: No changes this week — here's what's trending`

When zero changes occurred: use the trending/market line as the hook to avoid a dead email.

4.2 Full Email Template Structure

```
From: AIHackr Weekly <digest@aihackr.com>
Subject: Your AI Stack Digest: 3 changes across 12 watched domains this week

────────────────────────────────────────
           AIHackr Weekly Digest
           Week of Jan 13-20, 2025
────────────────────────────────────────

Your watchlist had 3 changes this week.
Anthropic continues to be the fastest-growing
provider across the SaaS tools you're tracking.

═══════════════════════════════════════
  CHANGES THIS WEEK  (3 of 12 domains)
═══════════════════════════════════════

1. linear.app
   ✅ ADDED Anthropic (Claude) — High confidence
   Detected: Jan 20 · Previously: OpenAI only
   → View full report

2. notion.so
   ❌ REMOVED OpenAI — High confidence (prior scan)
   ⬆️ UPGRADED Gemini: Medium → High confidence
   Detected: Jan 18
```

→ View full report

3. figma.com
   ⚠ CONFIDENCE CHANGE: Azure OpenAI High → Medium
   Note: May be a temporary detection fluctuation.
   Detected: Jan 16
   → View full report

YOUR WATCHLIST AT A GLANCE (12 domains)

Domain	Providers	Confidence	Last Changed
linear.app	Anthropic, OpenAI	High	Jan 20 ← NEW
notion.so	Gemini	High	Jan 18 ← NEW
figma.com	Azure OpenAI, OpenAI	Medium	Jan 16
vercel.com	OpenAI	High	Dec 12
stripe.com	OpenAI	High	Nov 3
loom.com	OpenAI, Anthropic	High	Oct 22
intercom.com	OpenAI	High	Oct 15
airtable.com	OpenAI, Azure OpenAI	High	Sep 30
miro.com	Gemini	Medium	Sep 12
clickup.com	OpenAI	High	Aug 28
coda.io	Anthropic	Medium	Aug 14
retool.com	None detected	—	—

[View Full Watchlist →]

TRENDING IN AI STACKS THIS WEEK

Across all domains monitored by AIHackr this week:

▲ Anthropic up +18%
  More SaaS tools added Claude this week than any other provider. Most common pattern: Anthropic added alongside (not replacing) an existing OpenAI integration — suggesting multi-model strategies.

▼ Standalone OpenAI down -6%
  OpenAI remains the most detected provider overall, but the share of tools using it exclusively continued to decline.

◆ Mistral spotted in 4 new products
  Early signal: Mistral is appearing in developer tools and European SaaS for the first time in our data.

◆ Multi-model stacks on the rise
  28% of monitored domains now show 2+ AI providers, up from 21% four weeks ago.

These trends are aggregated anonymously across

```
AIHackr's full domain database. Individual company
data is never shared.

═══════════════════════════════════════

[ Add More Domains to Watchlist → ]
[ Upgrade to Pro for Daily Scans → ]

──────────────────────────────────────
Sent to you@company.com because you have an AIHackr
watchlist. This digest sends every Monday at 8 AM
in your account timezone.

[ Manage digest settings ]  |  [ Unsubscribe ]
AIHackr · aihackr.com
──────────────────────────────────────
```

4.3 How the "Trending in AI Stacks" Section is Computed

The trending section is aggregated across all unique domains ever scanned by AIHackr (not limited to the user's watchlist), anonymized at the domain level. No user-identifiable data is exposed.

Data pipeline (runs Sunday night, results cache for Monday morning sends):

```
1. Pull all scan_results from the prior 7 days with scan_status = 'complete'
2. For each unique domain, take the most recent scan result of the week
3. Compare against the most recent scan result from the prior week for the same dom
4. Aggregate provider-level deltas:
   - provider_added count per provider
   - provider_removed count per provider
   - net_change = added - removed
   - pct_domains_using_provider (this week vs. last week)
5. Compute:
   - "Multi-provider" rate: domains with providers_detected.length >= 2
   - "New provider appearances": providers detected for the first time across the f
   - Fastest-growing provider: highest positive net_change
   - Notable movers: any provider with >2 standard deviations from its baseline wee
6. Write results to trending_weekly table (provider, week_start, added_count, remov
7. Template fills from trending_weekly WHERE week_start = current_week
```

Data shown in trending section:

- Maximum 4 trending bullets per digest
- Each bullet requires a minimum sample of 50 domain scans to be statistically meaningful (suppress if sample too small)
- Bullets are ranked by: (1) novelty — if a new provider is spotted for the first time, lead with that; (2) magnitude of change; (3) relevance to user's watchlist (if the user's watched domains are disproportionately affected, surface that angle)

5. Free vs Pro Gating

Feature	Free	Pro ($29/mo)
Watchlist size	Up to 5 domains	Unlimited
Re-scan cadence	Every 7 days	Every 24 hours
Manual "Scan Now"	1× per domain per week	3× per domain per day
Change detection history	Last 3 changes per domain	Full history (unlimited)
Email alerts	✅ On (up to 5 alerts/week total)	✅ On (unlimited)
Slack alerts	❌ Not available	✅ Available (webhook)
Alert threshold customization	Global only (no per-domain)	Global + per-domain overrides
Weekly digest	✅ Included (basic format)	✅ Included (full format with trending section)
Trending in AI Stacks section	❌ Not included in digest	✅ Included
Export watchlist data	❌ Not available	✅ CSV export
API access to watchlist data	❌ Not available	❌ (Enterprise only)
Watchlist sharing / team access	❌ Not available	❌ (Enterprise only, future)
Alert frequency cap	Max 1 alert/day per domain	Immediate (no cap)
Domain labels / notes	✅ Available	✅ Available
Pause monitoring	✅ Available	✅ Available
False positive dismissal	❌ Not available	✅ Available

Free plan design rationale:
The 5-domain limit is intentional — large enough to demonstrate value (a user with 5 competitors monitored is a retained user), small enough that any serious use case (sales prospecting, competitive tracking) immediately hits the ceiling. The 5-alerts/week total cap is the conversion trigger: once a user sees a watchlist change and gets the alert, the value is proven — but the cap interrupts the flow and creates urgency to upgrade.

6. Build Sequence

Phase 1: MVP (Weeks 1–4) — Ship the Loop

Week 1–2: Watchlist data model + manual add flow

- Implement `watchlist_entry` and `scan_result` tables
- Build "Add to Watchlist" button on scan result pages (single-domain add flow)
- Build basic Watchlist dashboard page with table view (no real-time updates yet)

- Trigger an immediate background scan on add
- **Why first:** Without the ability to save domains, nothing else in this spec exists. This is the keystone.

Week 2–3: Scheduled re-scan infrastructure

- Build scan queue with priority (Pro vs Free)
- Implement `next_scan_at` scheduling logic and cron job runner
- Implement `computeDiff` function and write to `stack_change` table
- **Why second:** Enables detection of changes — the core data for everything downstream.

Week 3–4: Email change alerts

- Build alert template and delivery via transactional email provider (Postmark / Resend)
- Implement deduplication logic (`alerted_at`, 24-hour cap)
- Build `/settings/alerts` with global email alert settings
- **Why third:** Email alerts are the simplest, highest-value signal delivery. They are also the proof-of-concept that the feature loop works end-to-end — add domain → change detected → user notified. This alone justifies Pro.

MVP deliverable: User can add 5 domains (free), get automatic re-scans weekly, receive an email when a provider changes. This is the complete retention loop, minimally viable.

Phase 2: Engagement Layer (Weeks 5–8)

Week 5–6: Weekly Digest

- Build digest aggregation pipeline (Sunday night job)
- Build digest email template (without "Trending" section — stub it with static copy for now)
- Implement digest delivery with user timezone awareness
- **Why next:** Digest creates a guaranteed weekly touchpoint regardless of whether changes happened. It also creates a habit of checking. Schedule reliability builds trust.

Week 6–7: Slack alerts

- Build webhook configuration UI in settings
- Implement Slack block kit payload delivery
- Per-domain Slack on/off override
- **Why here:** Slack delivery is the B2B stickiness multiplier. A Slack bot in a company's #competitive-intel channel creates team-level stickiness, not just user-level. Drives word of mouth.

Week 7–8: Alert settings (full)

- Per-domain alert threshold overrides
- Quiet hours

- Provider-specific filters
- False positive dismissal
- **Why here:** Reducing false-positive fatigue is what keeps users from unsubscribing from alerts. Ship this before alert volume grows large enough to create churn.

Phase 3: Intelligence Layer (Weeks 9–12)

Week 9–10: Trending in AI Stacks (digest section)

- Build `trending_weekly` aggregation pipeline
- Connect to digest template
- Build sample-size guardrails (minimum 50 domains)
- **Why now:** This is a Pro-only differentiator that justifies the subscription even if a user's watchlist is quiet. It also creates a reason for Pro users to stay even when nothing on their list changes.

Week 10–11: Export + change history

- Full change history view per domain (timeline UI)
- CSV export for Pro users
- **Why here:** Sales and competitive intelligence teams will want to show this data in slides. Export is a high-perceived-value, low-effort feature.

Week 11–12: Watchlist onboarding + plan gate UI

- Polish the upgrade prompt at the 5-domain limit
- Add "Upgrade for daily scans" contextual prompt when a change is detected on Free
- Email sequence: Day 1, Day 3, Day 7 post-watchlist-creation for users who haven't seen a change yet
- **Why last:** Conversion optimization belongs after the core loop is proven. You need real data on where users drop off before optimizing the funnel.

7. Success Metrics

The following KPIs measure whether the watchlist/alert loop is generating the intended retention and conversion behavior. All targets assume a 3-month post-launch evaluation window.

#	Metric	Definition	Target	Measurement method
1	Watchlist activation rate	% of registered users who add ≥1 domain to watchlist within 7 days of signup	≥40%	Event: `watchlist_entry.created` within 7 days of `user.created`

#	Metric	Definition	Target	Measurement method
2	Watchlist power user rate	% of active users with ≥5 domains on watchlist	≥ 20% (of activated users)	Count of users with watchlist_entry count ≥5
3	Week-4 retention (watchlist users)	% of users who added ≥1 watchlist domain and are still active (any session) at day 28	≥ 55% (vs. ≤30% for non-watchlist users)	Cohort retention analysis, segment by watchlist=yes/no
4	Alert open rate	% of change alert emails opened	≥ 45%	Email provider open tracking (Postmark/Resend)
5	Alert-to-upgrade conversion	% of free users who receive a change alert and upgrade to Pro within 14 days	≥ 8%	Funnel: `alert_sent` → `subscription.created` within 14 days
6	Digest open rate	% of weekly digest emails opened	≥ 35%	Email open tracking; measured per-send, not per-user
7	Free-to-Pro conversion (watchlist limit trigger)	% of free users who hit the 5-domain cap and upgrade within 7 days	≥ 15%	Funnel: `watchlist_limit_reached` event → `subscription.created` within 7 days

Secondary health metrics (monitor, no hard target):

- Alert unsubscribe rate (target: <2% per month — if higher, tune alert threshold defaults)
- Scan error rate per week (target: <5% of all scans in `error` state)
- False positive dismissal rate (target: <10% of all changes dismissed — if higher, improve diff algorithm)
- Slack webhook setup rate among Pro users (leading indicator of team stickiness)

Open Questions

#	Question	Owner	Blocking?
1	What is the expected scan infrastructure cost per domain per month? This drives the Free plan limit and Pro pricing floor.	Engineering	Yes — needed before launch
2	Should the weekly digest send day/time be user-configurable, or fixed Monday 8 AM (user TZ)? Fixed is simpler to build; configurable improves open rates.	PM	No — default Monday 8 AM for MVP
3	Do we need two consecutive missed-detection scans before recording PROVIDER_REMOVED, or is one scan sufficient? Lower sensitivity reduces false positives but delays real alerts.	PM + Eng	No — default 2 scans for MVP
4	Should the Trending section be powered by AIHackr's full scanned domain corpus or only domains on user watchlists? Full corpus gives better signal but requires explicit privacy policy language.	Legal + PM	No — use full corpus, add privacy note in email

#	Question	Owner	Blocking?
5	What transactional email provider? (Postmark preferred for deliverability; Resend for developer experience.)	Engineering	Yes — needed for Phase 1

End of spec — AIHackr Watchlist & Alerts v1.0

AIHackr Content & Leaderboard Engine — Full Playbook

Document: Part 3 — Content, SEO, and Leaderboard Strategy
Site: aihackr.com
Status: Ready for Monday execution
Last updated: 2025

EXECUTIVE SUMMARY

AIHackr's core SEO moat is a query category that no competitor owns: *"what AI does [company] use."* BuiltWith built a $30M+ business by owning "websites using [technology]" — AIHackr should replicate that model for the LLM era, layering a weekly Leaderboard on top to create a repeat-visit destination. The three-part engine: (1) programmatic company stack pages at scale, (2) a weekly AI Stack Leaderboard that drives links and shares, (3) blog content that ranks for informational queries and feeds both.

Zero-traffic sites do not climb from content quality alone. They climb from structural surface area: hundreds of pages each targeting a specific query, with a flagship asset that earns ongoing backlinks. This document is that architecture.

SECTION 1: PROGRAMMATIC SEO — COMPANY STACK PAGES

1.1 URL Structure

Recommended format:

```
aihackr.com/stack/[company-slug]
```

Examples:

```
aihackr.com/stack/notion
aihackr.com/stack/linear
aihackr.com/stack/salesforce
aihackr.com/stack/hubspot
aihackr.com/stack/intercom
```

Rationale: `/stack/` is short, memorable, and functions as both a noun and verb ("check the stack"). Avoid using domain-based slugs (e.g., `/stack/notion.so`) as the company name is what people search for, not the domain. Use the most-recognized short name, all lowercase, hyphenated for multi-word names (e.g., `/stack/harvey-ai`, `/stack/adobe-firefly`).

Canonical handling: Each page should set a canonical to itself. If the same company has two slugs (e.g., `/stack/openai` and `/stack/open-ai`), 301 redirect the variant to the canonical.

Pagination for future scale: When any company has 5+ historical scan entries, add `/stack/[slug]/history` sub-page. Do not paginate the main page — keep it as one scrollable page.

1.2 Full Page Template

Title Tag Formula

```
[Company Name] AI Stack — Which LLM Does [Company Name] Use? | AIHackr
```

Examples:

- `Notion AI Stack — Which LLM Does Notion Use? | AIHackr`
- `Intercom AI Stack — Which LLM Does Intercom Use? | AIHackr`
- `Salesforce AI Stack — Which LLM Does Salesforce Use? | AIHackr`

Rules: Keep under 60 characters before `| AIHackr`. If the company name is long, use: `[Company] AI Provider — AI Stack Check | AIHackr`.

Meta Description Formula

```
See which AI provider [Company Name] uses in production. AIHackr scanned [Company N
```

Example:

```
See which AI provider Notion uses in production. AIHackr scanned Notion on June 2 ar
```

H1 Formula

```
[Company Name]'s AI Stack: Which LLM Are They Using?
```

Examples:

- `Notion's AI Stack: Which LLM Are They Using?`
- `Salesforce's AI Stack: Which LLM Are They Using?`

Intro Paragraph (2 sentences, with placeholders)

> [COMPANY] has integrated AI capabilities across its product — but which large language model (LLM) is actually running under the hood? AIHackr scanned [COMPANY]'s public-facing endpoints and detected [PRIMARY_PROVIDER] as the primary AI provider, with signals including [TOP_EVIDENCE_SIGNAL_1] and [TOP_EVIDENCE_SIGNAL_2].

Variant when confidence is low (< 60%):

> [COMPANY] appears to use AI in several product features, though the specific LLM provider is difficult to confirm from external signals alone. AIHackr's latest scan found [SIGNAL_DESCRIPTION], which may indicate [PROVIDER], but [COMPANY] may be using a private or self-hosted model.

Data Card — Fields Ordered by Priority

Display as a structured card (grid layout, 2 columns on desktop, 1 on mobile). Fields in this exact order:

Priority	Field	Display Label	Notes
1	Primary AI Provider	AI Provider	Name + logo chip (e.g., "Anthropic Claude")
2	Confidence Score	Detection Confidence	Percentage + color badge (green ≥ 80%, yellow 50–79%, red < 50%)
3	Last Scan Date	Last Scanned	Relative + absolute (e.g., "3 days ago · June 2, 2025")
4	Evidence Signals	Evidence	Count of signals with expand link ("4 signals detected — see below")
5	AI Framework	AI Framework	e.g., LangChain, LlamaIndex, custom, unknown
6	Hosting / Inference	Inference Host	e.g., AWS Bedrock, Azure OpenAI, direct API, self-hosted
7	Secondary Provider	Secondary AI Provider	Show only if detected; e.g., "OpenAI (embeddings)"
8	Model Version	Model Version	e.g., "Claude 3.5 Sonnet (inferred)" — label as inferred if not confirmed
9	Stack Category	Company Category	e.g., "CRM", "Dev Tools", "Support"
10	Scan History	History	"First detected [date] · [N] scans"

Below the card: Two CTA buttons side by side:

- Primary: `Track [Company] Changes` (subscription/alert)
- Secondary: `Export Stack Data` (requires account)

"Evidence" Section

Section heading: `How We Detected This`

Sub-heading: `Signals found during the [DATE] scan`

Display evidence as a scannable list with signal type icons. Each entry shows:

```
[ICON] [Signal Type]: [Signal Description]
      Confidence contribution: [HIGH / MEDIUM / LOW]
      Detected at: [URL or endpoint, abbreviated]
```

Signal types to display (with icons):

Signal Type	Icon	Example
HTTP Response Header	📡	`x-llm-provider: anthropic` found in `/api/chat` response
JavaScript SDK Reference	📦	`@anthropic-ai/sdk` loaded on `app.notion.so`
Network Request Pattern	🌐	Outbound request to `api.anthropic.com` observed
Streaming Response Format	⚡	Server-sent events format matches Claude streaming API
Error Message Fingerprint	🔍	Error string `claude-3-sonnet` found in client-side JS
DNS / IP Resolution	🗺	Requests routing through AWS Bedrock (inferred Anthropic)
Model Response Signature	🧬	Response token patterns match Claude model family
Public Disclosure	📢	Company blog post or changelog references [Provider]

Below evidence list: Small disclaimer in muted text:

> Detection is based on observable external signals only. AIHackr does not access private APIs or internal systems. Confidence scores reflect signal strength, not certainty.

"Historical Changes" Section

Section heading: `AI Stack History`

Sub-heading: `How [Company]'s AI provider has changed over time`

Display as a vertical timeline (newest at top):

```
● [DATE] — Current
  Provider: Anthropic (Claude 3.5 Sonnet)
  Confidence: 89%
  Change: Upgraded from Claude 3 Opus

● [DATE] — 3 months ago
  Provider: Anthropic (Claude 3 Opus)
  Confidence: 82%
```

```
  Change: Switched from OpenAI GPT-4

● [DATE] — 8 months ago
  Provider: OpenAI (GPT-4)
  Confidence: 75%
  Change: First detected
```

When no history exists (new page): Show a single entry with `Scan Date: [DATE] — First scan. Check back for changes.`

Below timeline: CTA: `Get notified when [Company]'s AI stack changes →` (email alert sign-up, inline, no redirect)

"Similar Companies Using [Same AI Provider]" Section

Section heading: `Other Companies Using [PROVIDER]`

Sub-heading: `[N] companies in our database also use [PROVIDER] — here are the most-watched:`

Display as a 3-column card grid (2 on tablet, 1 on mobile), maximum 6 cards. Each card:

```
[Company Logo]
[Company Name]
[Category Badge]
[Confidence %]
→ View AI Stack
```

Sorting logic for which 6 to show: Prioritize companies in the same category as the current company first, then by confidence score descending. Do not show the current company.

Link: `See all [N] companies using [PROVIDER] →` — links to `/provider/[provider-slug]` (e.g., `/provider/anthropic`)

Schema Markup

Use **TechArticle** schema as the primary type, nested inside a **WebPage**. Do not use FAQ (thin justification) or SoftwareApplication (wrong entity type).

```
{
  "@context": "https://schema.org",
  "@type": "TechArticle",
  "headline": "[Company Name]'s AI Stack: Which LLM Are They Using?",
  "description": "Analysis of [Company Name]'s AI provider based on external signal
  "datePublished": "[FIRST_SCAN_DATE]",
  "dateModified": "[LAST_SCAN_DATE]",
  "author": {
    "@type": "Organization",
    "name": "AIHackr",
```

```
    "url": "https://aihackr.com"
  },
  "publisher": {
    "@type": "Organization",
    "name": "AIHackr",
    "logo": {
      "@type": "ImageObject",
      "url": "https://aihackr.com/logo.png"
    }
  },
  "about": {
    "@type": "SoftwareApplication",
    "name": "[Company Name]",
    "url": "https://[company-domain]",
    "applicationCategory": "[Category]"
  },
  "mentions": {
    "@type": "SoftwareApplication",
    "name": "[AI Provider]",
    "url": "https://[provider-url]"
  }
}
```

Also add **BreadcrumbList** schema:

```
{
  "@type": "BreadcrumbList",
  "itemListElement": [
    {"@type": "ListItem", "position": 1, "name": "Home", "item": "https://aihackr.
    {"@type": "ListItem", "position": 2, "name": "AI Stack Explorer", "item": "htt
    {"@type": "ListItem", "position": 3, "name": "[Company Name]", "item": "https:
  ]
}
```

CTA Placement and Copy

CTA 1 — Inline, immediately after the data card (above the fold on desktop):

- Button: `Track [Company]` (primary, filled)
- Sub-text: `Get an email when their AI stack changes. Free.`

CTA 2 — After the Evidence section:

- Button: `Run a Fresh Scan` (secondary, outlined)
- Sub-text: `Last scanned [DATE]. Request an updated scan.`

CTA 3 — After the Historical Changes section:

- Banner: `Are you from [Company]? Claim this page to add context or correct detections.`
- Link: `Claim this page →`

CTA 4 — Bottom of page (exit intent / scroll depth):

- Box with headline: `What AI does your SaaS use?`
- Sub: `Scan any domain. Free for the first 5 companies.`
- Button: `Try the Scanner →` (links to homepage scanner)

1.3 Seed List — 50 High-Value SaaS Companies

Prioritized by: name recognition, high search volume, confirmed AI feature presence, likelihood of detectable provider signals.

Tier 1 — Consumer-facing AI products (highest search volume):

1. Notion
2. Intercom
3. Zendesk
4. HubSpot
5. Salesforce (Einstein AI)
6. Linear
7. Figma (AI features)
8. Canva
9. Grammarly
10. Duolingo

Tier 2 — AI-native or heavily AI-integrated SaaS:

11. Cursor
12. GitHub Copilot (Microsoft)
13. Jasper
14. Copy.ai
15. Writesonic
16. Perplexity
17. Harvey (Harvey.ai)
18. Casetext
19. Glean
20. Guru

Tier 3 — Enterprise SaaS with AI additions:

21. Workday
22. ServiceNow
23. Freshdesk
24. Drift (Salesloft)
25. Gong
26. Chorus (ZoomInfo)
27. Outreach

28. Salesloft
29. Apollo.io
30. Clay

Tier 4 — Developer and infrastructure tools:
31. Vercel (v0.dev)
32. Replit
33. Codeium
34. Tabnine
35. Sourcegraph (Cody)
36. JetBrains AI
37. Sentry AI
38. Datadog AI
39. New Relic (applied intelligence)
40. Cloudflare AI

Tier 5 — High-profile companies with known AI bets:
41. Brex
42. Ramp
43. Stripe (Stripe Radar / AI)
44. Plaid
45. Rippling
46. Deel
47. Lattice
48. Leapsome
49. Typeform
50. Loom (Atlassian)

Rationale for this order: Tier 1 names get searched by name millions of times per month; even a small fraction of "[Company] AI" queries converts to page visits. Tier 2 companies are already AI-forward, so readers searching "what AI does Jasper use" have very specific intent. Tiers 3–5 add long-tail coverage while expanding topical authority across categories.

1.4 How to Scale — Automation Strategy

Phase 1: Manual scaffold, automated data (Weeks 1–4)

1. **Build the template once** in your CMS or static site generator (Next.js recommended). The template reads from a JSON/database record — no page-specific code.
2. **Create seed database** with 50 records: `{slug, company_name, category, domain, primary_provider, confidence, evidence[], last_scanned, history[]}`.
3. **Publish all 50 pages** in one batch. Date-modified should be your actual scan date, not launch date.
4. **Submit sitemap** to Google Search Console immediately. Index all 50 pages manually using the URL Inspection tool.

Phase 2: Scan automation (Weeks 4–8)

1. **Scheduled scanner:** Run AIHackr's detection engine on every company in the DB on a 7-day cadence. Write results to DB. Template auto-updates the page on next build.
2. **Change detection:** If `primary_provider` or `confidence` changes ≥ 10 points vs. last scan, flag for human review before publishing. This prevents noisy automated corrections.
3. **Trigger rebuild:** Use a webhook from DB → build pipeline (Vercel, Netlify, or custom) to regenerate only updated pages. Full-site rebuild is unnecessary.

Phase 3: Discovery and expansion (Weeks 8–16)

1. **Search query mining:** Use Google Search Console data weekly. Any query containing "[company name] AI" or "what AI does [company]" that (a) gets impressions and (b) has no corresponding page → add that company to the queue.
2. **Backlink gap:** Use Ahrefs or Semrush to find which company names drive inbound links to competitors' tech-stack pages. Add those companies.
3. **User requests:** Add a `request a company` form on the main `/stack` index page. Users self-identify the next 100 companies.

Phase 4: Scale to 1,000+ pages (Month 4+)

1. **Batch import:** Scrape (from public sources like Crunchbase, G2, ProductHunt) a list of SaaS companies with known AI features. Auto-create stub pages with minimal data ("scan not yet run") and schedule scans.
2. **Stub pages:** Even pages with `confidence: null` and `status: pending_scan` have SEO value — they will appear in search results as "AIHackr has not yet scanned this company. Request a scan."
3. **Minimum human touch per page:** ~3 minutes — one human reviews the auto-generated confidence rating and evidence list before the page goes live. After 30 days of accurate results, skip human review for low-risk updates (provider unchanged, confidence change < 5 points).

Update cadence

- Re-scan every company every 7 days
- Auto-publish non-material changes (confidence drift < 10%)
- Queue material changes (provider switch, new framework detected) for 24-hour human review
- Mark each page with a `data-freshness` indicator so users know when it was last checked

SECTION 2: THE AI STACK LEADERBOARD

2.1 Concept Summary

A public, shareable page showing the AI provider distribution across a curated set of companies — updated weekly, with a "what changed this week" panel that gives people a reason to return and share. The leaderboard's primary value is not the full table (though it has one) — it's the weekly delta: who switched, who upgraded, who is hedging with a second model.

2.2 Page Design Spec

URL

```
aihackr.com/leaderboard
```

Permanent URL — do not date-stamp it. Weekly state is maintained on the page; individual weekly snapshots can live at `/leaderboard/2025-W23` for archival/sharing.

Headline and Subheadline Copy

Headline:

```
The AI Stack Leaderboard
```

Subheadline:

```
Which LLMs are powering the world's top SaaS companies — updated every Monday.
```

Above-the-fold description (2 sentences):

```
AIHackr scans hundreds of SaaS products weekly to detect which AI providers they use
This leaderboard tracks the AI stack of companies across YC cohorts, B2B SaaS, and d
and surfaces what changed since last week.
```

Freshness badge (always visible):

```
Last updated: [DATE] · Next update: [DATE] · [N] companies tracked
```

Main Data Table — Columns

Column	Display Name	Sortable	Notes
—	Rank	No	Auto-number based on current sort
—	Company	Yes	Logo + name + category badge
—	Primary AI Provider	Yes	Provider logo chip + name
—	Confidence	Yes	% score + color badge
—	AI Framework	Yes	LangChain, LlamaIndex, custom, unknown
—	Inference Host	Yes	AWS Bedrock, Azure OpenAI, Direct API, Self-hosted
—	Secondary Provider	No	Shown as small tag if detected, blank if not
—	Last Scanned	Yes	Relative date
—	Week-over-Week	No	"↔ No change", "↑ New", "⊠ Switched from [X]"
—	Watch	No	Star/bell icon to add to watchlist

Default sort: Last Scanned (newest first), then alphabetical by company name.

Available sort options:

- Most recently scanned
- Highest confidence
- Alphabetical (A–Z)
- By AI provider (groups companies by provider)
- By category

Filters

- **AI Provider:** OpenAI · Anthropic · Google Gemini · Mistral · Meta Llama · Cohere · AWS Bedrock · Azure OpenAI · Other · Unknown
- **Category:** CRM · Dev Tools · Support · Productivity · Finance · HR · Marketing · Data · Security · Other
- **Confidence:** High (≥80%) · Medium (50–79%) · Low (<50%)
- **YC Batch:** W24 · S24 · W25 · S25 · W26 · Non-YC
- **Changed this week:** Toggle (show only companies with new detections or provider changes)

"This Week's Changes" Panel

Position: Pinned to top right on desktop, collapsed to an expandable drawer on mobile.

Format:

```
━━━━━━━━━━━━━━━━━━━━
□ WEEK OF [DATE]
━━━━━━━━━━━━━━━━━━━━

□ PROVIDER SWITCHES (N)
• [Company] switched from OpenAI → Anthropic
• [Company] added Gemini as secondary provider

⬆ NEW ADDITIONS (N)
• [Company] — first detected: OpenAI GPT-4o
• [Company] — first detected: Anthropic Claude

□ CONFIDENCE CHANGES (N)
• [Company] confidence ↑ 68% → 84% (stronger signal)
• [Company] confidence ↓ 79% → 61% (signals weakened)

□ THIS WEEK'S STAT
OpenAI leads with [N] companies ([X]%)
Anthropic holds [N] companies ([X]%)
━━━━━━━━━━━━━━━━━━━━
Share this week's update → [icon row]
```

The panel should be sticky on desktop scroll. On mobile, show as a dismissable banner.

Embed / Share Feature

1. Shareable badge (per company):

```
┌──────────────────────────┐
│ □ [Company] uses         │
│    Anthropic Claude      │
│    87% confidence        │
│    Powered by AIHackr    │
└──────────────────────────┘
```

SVG badge, auto-generated at `/badge/[slug].svg`. Companies can embed in their own README or website (though most won't — this is for community/developer audiences).

2. Leaderboard embed (iframe):

Embeddable version of the full leaderboard at:

```
https://aihackr.com/embed/leaderboard?filter=yc-w25&theme=light
```

Query params: `filter`, `theme` (light/dark), `limit` (default 10, max 50), `provider`.

Marketing blogs, newsletters, and VC portfolio pages can embed a live version.

3. Weekly snapshot image:

Auto-generated weekly PNG/OG image (1200×630px) with:

- "AI Stack Leaderboard — Week of [DATE]"
- Top 10 companies with provider chips
- Provider distribution pie chart
- AIHackr branding

Accessible at: `/leaderboard/2025-W23/og.png`

This image is what populates Twitter/X cards, LinkedIn thumbnails, and newsletter embeds without any extra work.

4. Direct table link with pre-filtered state:

Every filter combination generates a shareable URL:

```
aihackr.com/leaderboard?provider=anthropic&category=devtools
```

The URL state is preserved so links shared in Slack/Discord/Twitter show exactly what the sharer saw.

How It Updates

Fully automated with a human "publish gate":

1. Sunday night: scanner runs on all tracked companies, writes results to staging DB.
2. Monday 6am: automated diff script generates the "This week's changes" content and creates draft publish.
3. Monday 7am: human reviews the diff (5–10 minutes): checks for false positives in provider switches, verifies any unusual confidence drops, approves publish.
4. Monday 8am: publish. Simultaneously triggers the Twitter/X thread draft (see Section 2.4), sends newsletter content block to the email queue, and posts the OG image to the scheduled social queue.

Total human time per weekly update: 20–30 minutes (review + approve + minor copy edits to social drafts).

2.3 Leaderboard Segments at Launch

Segment A: Top 100 YC AI Companies (W24–W26)

Description: The 100 highest-profile YC companies from recent batches that have shipped AI features. Source: YC company directory filtered to AI/ML + B2B SaaS.

Why this segment works: YC founders follow each other obsessively. A leaderboard of YC AI companies gets shared inside YC Slack, founder forums, and VC portfolio updates. Every company in the list has an incentive to share their own entry.

Segment B: Top AI-Native B2B SaaS (20 companies to launch with)

The following 20 are verified AI-native B2B SaaS with high likelihood of detectable providers:

1. **Cursor** — AI code editor (OpenAI + Anthropic dual-provider likely)
2. **Codeium** — AI coding assistant
3. **Sourcegraph Cody** — AI code search and generation
4. **Harvey** — AI legal research
5. **Casetext** — AI legal drafting
6. **Glean** — enterprise AI search
7. **Guru** — AI knowledge management
8. **Jasper** — AI marketing copy
9. **Writer** — enterprise AI writing
10. **Cohere** — AI for enterprise (also an AI provider itself)
11. Otter.ai — AI meeting notes
12. Fireflies.ai — AI meeting assistant
13. **Gong** — AI revenue intelligence
14. **Chorus** — AI conversation intelligence
15. **Moveworks** — AI IT helpdesk
16. **Aisera** — AI service management
17. **Forethought** — AI customer support
18. **Kustomer** — AI CRM
19. **Rasa** — AI conversational AI (open-source adjacent)
20. **Cohere for Enterprise** (distinct from Cohere as provider)

Segment C: Most-Watched Companies

Description: Dynamically generated — the top 50 companies by number of users who have starred/watched them in AIHackr.

This segment is the most powerful for retention: users see their watchlist companies aggregated on a public leaderboard, which reinforces the value of creating an account and watching companies.

At launch (before critical watchlist mass): Seed this segment with the 50 companies from Section 1.3 (the seed list). As real watchlist data accumulates (target: 200+ registered users), transition to genuine watchlist data.

2.4 Distribution Plan for Each Weekly Drop

Twitter/X Thread Template

Tweet 1 (Hook — this is the one that gets retweeted):

```
The AI Stack Leaderboard is live for week of [DATE].

[N] companies tracked. [N] changes this week.

The biggest move: [Company] switched from [OLD PROVIDER] to [NEW PROVIDER].

Here's what else shifted 
```

Tweet 2 (Provider distribution stat):

```
This week's AI provider breakdown across [N] top SaaS companies:

OpenAI: [N] companies ([X]%)
Anthropic: [N] companies ([X]%)
Google Gemini: [N] companies ([X]%)
Self-hosted/OSS: [N] companies ([X]%)

Anthropic is [gaining/losing] [X] points week-over-week.
```

Tweet 3 (Biggest mover):

```
Most interesting move this week:

[Company] → [NEW PROVIDER]

What we detected: [1-2 sentence description of evidence type]
Confidence: [X]%

Why it matters: [1-2 sentence analysis of what this switch signals for the market]
```

Tweet 4 (New additions):

```
[N] companies added to the leaderboard this week:

• [Company] — [Provider] ([X]% confidence)
• [Company] — [Provider] ([X]% confidence)
• [Company] — [Provider] ([X]% confidence)

Full list: aihackr.com/leaderboard
```

Tweet 5 (Category angle):

```
Among [Category] tools specifically:

[X]% are on Anthropic
[X]% are on OpenAI

The pattern: [1-sentence insight about why this category skews toward a particular pr
```

Tweet 6 (CTA):

```
The full AI Stack Leaderboard is at aihackr.com/leaderboard

Free to view. Subscribe to get next Monday's update in your inbox.

Which company switch surprised you most?
```

Posting time: Monday 9am PT. Thread format (all 6 tweets in sequence). Quote-tweet the thread from the @AIHackr account with a one-liner summary for maximum algorithmic reach.

HN Submission Strategy

Submit on Monday afternoons (12pm–3pm PT) — HN has highest active users during US business hours.

Title formula (choose based on what the week's biggest story is):

Option A (provider switch story): `[Company] switched from [Old Provider] to [New Provider] — we detected it before they announced it`

Option B (data/trend story): `OpenAI vs. Anthropic in production: our data across [N] SaaS companies`

Option C (tool announcement): `Show HN: AI Stack Leaderboard — we scan [N] SaaS companies weekly to see which LLM they use`

Rules:

- Never submit the same URL twice. Each week must link to the dated snapshot (e.g., `/leaderboard/2025-W23`) not the root leaderboard.
- The first comment should be from the founder: a brief explanation of the methodology and a genuine "ask HN" question to seed discussion (e.g., "Curious if anyone can explain why nearly all YC W25 companies are on Anthropic — is this a partnership deal or just model preference?")
- Do not submit on the same week as a major OpenAI or Anthropic product launch — those threads will bury you.

LinkedIn Post Format

Format: Single-image post (use the weekly OG image) with copy:

```
The AI Stack Leaderboard — Week of [DATE].

[N] companies. [N] changes.

The headline this week: [1-sentence biggest story].

What this tells us: [3-sentence market insight — who's winning the enterprise AI race

Full breakdown (free): [link]

Three things I found interesting this week:

1. [Observation about a specific company or category]
2. [Observation about provider trends]
3. [Something counterintuitive or surprising]

Which AI provider do you think wins enterprise in 2026?
```

Target accounts to tag (where relevant): Tag the company(s) involved in the biggest story — they often comment or share, which boosts reach significantly. Tag the AI provider if they announced something that week.

Newsletter Pitch Angle

Target newsletters (pitch for inclusion or content partnership):

Newsletter	Audience	Pitch angle
TLDR AI	500k+ tech/AI readers	Weekly data item: "X companies switched providers this week"
The Batch (deeplearning.ai)	ML practitioners	"Real production LLM adoption data from [N] SaaS companies"
Ben's Bites	AI founders/investors	"Which AI provider is winning among top SaaS?"
The Pragmatic Engineer	Senior engineers	"What LLMs are SaaS companies actually shipping"

Newsletter	Audience	Pitch angle
Lenny's Newsletter	PMs, product leaders	"How to pick your LLM: data from [N] production apps"
SaaStr Weekly	SaaS founders/investors	"Which AI provider do the fastest-growing SaaS companies use?"
NFX Signal	VC/founders	"Network effects data: why Anthropic is winning YC companies"

Pitch structure (for cold outreach):

- Subject: `Weekly data: which AI provider [Newsletter's audience] actually use`
- Body: 3 sentences max. Attach the weekly OG image. Offer exclusive early access to the leaderboard data for their newsletter readers (a filter or segment they can't get elsewhere).

Why This Is Link-Bait — Why VCs and PMs Share It

For VCs:

- Portfolio monitoring: VCs want to know if their portfolio companies are using expensive inference APIs (cost implication) or have switched to a competitor's model (strategic signal). The leaderboard is a free intelligence feed.
- Due diligence signal: "Which LLM does [target company] use?" is a real due diligence question. AIHackr answers it publicly, and VCs share the link when it answers a question they've already been asking.
- Bragging rights: If a VC's portfolio company is #1 on the leaderboard, they tweet about it.

For PMs and engineers:

- Social proof: "Who else is using X model?" is a constant question in engineering discussions. The leaderboard answers it with receipts.
- Competitive intelligence: Knowing that a competitor switched providers is a legitimate strategic input. PMs share this to their team Slacks immediately.
- The "aha" moment: When someone finds their own company on the leaderboard (correctly or incorrectly detected), they share it — either to validate or dispute. Both generate traffic.

The shareable moment in every weekly drop: Always engineer one "wow" finding per week. Even if the data is thin, find the most counterintuitive data point (e.g., "The #1 CRM switched to Gemini — OpenAI's enterprise moat may be weakening") and make that the lead. This is the hook that gets screenshotted and quoted.

SECTION 3: BLOG CONTENT CALENDAR — 90 DAYS (12 WEEKS)

Publishing cadence: 1 post per week (12 posts total). Posts publish Tuesday at 9am (one day after the leaderboard drops Monday, so social traffic is fresh).

Word count guidance: Target lengths are floors, not ceilings. Do not pad. A 1,200-word post with original data beats a 2,500-word post with generic content.

Week 1

Title: `What AI Does Notion Use? A Deep Dive Into Notion AI's LLM Stack`
Target keywords: "what AI does Notion use", "Notion AI LLM", "Notion AI provider"
Search intent: Informational / product research (users wondering what powers Notion AI)
Word count: 1,400
Key angle: First-party detection data from AIHackr's scan. Explain what signals were found, what they indicate about provider choice (Anthropic Claude), and why Notion likely made that choice. Include a comparison to when they were previously on OpenAI (if applicable). Surface the confidence score prominently.
Internal linking: Links to `/stack/notion`, `/leaderboard`, and the homepage scanner

Week 2

Title: `OpenAI vs. Anthropic in Production: Data From 50 SaaS Companies`
Target keywords: "OpenAI vs Anthropic production", "which LLM do companies use", "AI provider comparison enterprise"
Search intent: Research / comparative (engineers and PMs evaluating LLM choices)
Word count: 2,200
Key angle: Original data post using AIHackr scan results across the 50-company seed list. Show provider distribution, confidence averages by provider, and category breakdown (e.g., "Support tools skew Anthropic; Dev tools skew OpenAI"). This is the cornerstone data post — every future comparative post references it.
Internal linking: Links to `/leaderboard`, 5 individual `/stack/` pages referenced in the data, `/stack` index

Week 3

Title: `What AI Does Intercom Use? Detecting the LLM Behind Fin AI`
Target keywords: "what AI does Intercom use", "Intercom Fin AI LLM", "Intercom AI provider"
Search intent: Informational / product research
Word count: 1,200
Key angle: Intercom's Fin AI is one of the most-watched AI support products. Break down what AIHackr detected, note any evidence around their model choice, and contextualize against other support-tool AI providers (Zendesk, Freshdesk). Cross-link to those pages.
Internal linking: Links to `/stack/intercom`, `/stack/zendesk`, `/stack/freshdesk`, `/leaderboard`

Week 4

Title: `How to Detect Which LLM a Website Is Using: A Technical Guide`
Target keywords: "how to detect which LLM a website uses", "detect AI provider website", "LLM fingerprinting"
Search intent: Informational / how-to (developers and researchers)
Word count: 2,500
Key angle: Explain the detection methodology AIHackr uses — HTTP headers, JS SDK signatures, network request patterns, streaming response formats, error fingerprinting. This is the "methodology" post that gives credibility to all other detection claims. Positions AIHackr as the authoritative source, not just a directory. Include code examples (Python/curl) for readers who want to try it themselves.
Internal linking: Links to homepage scanner, `/stack` index, Week 2 data post

Week 5

Title: `YC W25 AI Companies: Which LLM Providers Are They Actually Using?`
Target keywords: "YC W25 AI companies LLM", "YC companies using OpenAI", "YC companies using Anthropic"
Search intent: Research / leaderboard-adjacent
Word count: 1,800
Key angle: Original data using AIHackr scans of detected YC W25 companies. Show provider distribution table, highlight any companies that switched mid-cohort, note which companies have undetectable stacks (likely self-hosted or custom). Tag this post with a "Updated monthly" badge.
Internal linking: Links to `/leaderboard` (YC segment), 3–5 individual `/stack/` pages, Week 2 data post

Week 6

Title: `What AI Does HubSpot Use? Inside HubSpot's AI Stack`
Target keywords: "what AI does HubSpot use", "HubSpot AI provider", "HubSpot LLM"
Search intent: Informational / product research
Word count: 1,300
Key angle: HubSpot AI features are well-known (Content Hub, Breeze AI). Detection angle: what does the scan reveal about whether they're using their own fine-tuned model, an OpenAI partnership, or another provider? If HubSpot has made public announcements, weave that in as "confirmed signal" alongside the technical detection.
Internal linking: Links to `/stack/hubspot`, `/stack/salesforce` (competitor), `/leaderboard`

Week 7

Title: `The Rise of Anthropic in Enterprise SaaS: Data From 100 Companies`
Target keywords: "Anthropic enterprise adoption", "Claude enterprise SaaS", "Anthropic vs OpenAI enterprise"
Search intent: Research / trend (investors, PMs, market analysts)
Word count: 2,000
Key angle: Second major data post. Use AIHackr scan data to chart Anthropic's market share growth across the tracked company list over time (even 8 weeks of data shows a trend). Compare enterprise categories — legal tech, support, productivity — to show where Anthropic is winning. This post specifically targets VC/investor readers who will link from their own newsletters.
Internal linking: Links to `/leaderboard`, `/provider/anthropic` (planned page), Week 2 data post

Week 8

Title: `What AI Does Salesforce Use? Einstein GPT and the LLM Behind the CRM`
Target keywords: "what AI does Salesforce use", "Salesforce Einstein AI provider", "Salesforce LLM"
Search intent: Informational / product research
Word count: 1,400
Key angle: Salesforce is unique — they have a public partnership with OpenAI and also use multiple providers. Detection complexity is high. This post can productively show the limits of external detection (some signals found, confidence medium) while still being genuinely useful. Honesty about confidence = credibility.
Internal linking: Links to `/stack/salesforce`, `/stack/hubspot`, `/leaderboard`

Week 9

Title: `AI Stack Checker Tools: How to See Which LLM Any SaaS Is Using`
Target keywords: "AI stack checker", "detect AI provider tool", "LLM stack detection tool"
Search intent: Tool discovery / transactional
Word count: 1,600
Key angle: Comparison post that naturally positions AIHackr as the best tool. Compare the manual detection approaches (DIY with curl + JS inspection) vs. AIHackr's automated approach. Include a brief mention of what BuiltWith does (tech stack detection) and why that's different from AI-specific detection. Avoid being sycophantic — the post should be genuinely useful even if someone never uses AIHackr.
Internal linking: Links to homepage scanner, `/stack` index, Week 4 how-to post

Week 10

Title: `What AI Does Cursor Use? Inside the AI Code Editor's LLM Stack`
Target keywords: "what AI does Cursor use", "Cursor LLM provider", "Cursor AI model"
Search intent: Informational / high-search-volume (Cursor is extremely popular with developers)
Word count: 1,500
Key angle: Cursor is known to use both OpenAI and Anthropic. Detection angle: which model family do they default to? What does the scan evidence show about which provider handles which feature (autocomplete vs. chat vs. code generation)? Developer-audience post, so technical depth is appropriate.
Internal linking: Links to `/stack/cursor`, `/stack/codeium`, `/stack/sourcegraph-cody`, `/leaderboard`

Week 11

Title: `Why Companies Switch LLM Providers: Patterns From 3 Months of Scan Data`
Target keywords: "why companies switch LLM provider", "OpenAI to Anthropic switch", "LLM switching patterns enterprise"
Search intent: Research / analytical (engineers, PMs evaluating their own provider choices)
Word count: 2,000
Key angle: Third major data post — now with ~3 months of scan history. Show the actual provider switches detected in the AIHackr database. Categorize switch reasons (inferred, not stated): cost, capability, privacy, enterprise contracts. This post generates the most links because it answers a question everyone in AI product is asking. It is also highly shareable on LinkedIn among AI engineers.
Internal linking: Links to `/leaderboard`, Week 2 and Week 7 data posts, `/stack` pages of companies referenced

Week 12

Title: `The AI Stack Report: Q3 2025 — Which LLMs Are SaaS Companies Deploying?`
Target keywords: "AI stack report 2025", "LLM adoption enterprise 2025", "which AI providers SaaS 2025"
Search intent: Research / report (investors, analysts, enterprise buyers)
Word count: 3,000
Key angle: The quarterly report — a comprehensive synthesis of all scan data collected in the first 90 days. Provider market share, trend lines, category breakdowns, notable switches, confidence distributions. This is the link magnet post: designed to be cited by other publications, newsletters, and VC memos. Include downloadable summary PDF. Pitch this to TLDR AI, Ben's Bites, and SaaStr simultaneously with the publish date.
Internal linking: Links to `/leaderboard`, all prior data posts (Weeks 2, 7, 11), `/stack` index, 10+ individual company pages

SECTION 4: KEYWORD TARGETS

Cluster 1: Brand-Adjacent — "What AI Does [X] Use?" (10 queries)

Keyword	Search Intent	Difficulty	Target Page
what AI does Notion use	Informational	Medium	`/stack/notion` + Week 1 blog post
what AI does Intercom use	Informational	Low	`/stack/intercom` + Week 3 blog post
what AI does HubSpot use	Informational	Medium	`/stack/hubspot` + Week 6 blog post
what AI does Salesforce use	Informational	Medium	`/stack/salesforce` + Week 8 blog post
what AI does Cursor use	Informational	Low–Medium	`/stack/cursor` + Week 10 blog post
what AI does Zendesk use	Informational	Low	`/stack/zendesk`
what AI does Linear use	Informational	Low	`/stack/linear`
what AI does Grammarly use	Informational	Low	`/stack/grammarly`
what AI does Duolingo use	Informational	Low	`/stack/duolingo`
what AI does Gong use	Informational	Low	`/stack/gong`

Notes: These are long-tail queries with very specific intent. Difficulty is low on most because no authoritative page currently answers them. AIHackr can own this entire cluster within 3 months of publishing the pages.

Cluster 2: Category — "Detect AI Provider" / "AI Stack Checker" (5 queries)

Keyword	Search Intent	Difficulty	Target Page
detect AI provider website	Tool discovery	Low	Homepage / Week 9 blog post
AI stack checker	Tool discovery	Low	Homepage
LLM stack detection tool	Tool discovery	Low	Homepage / Week 4 blog post
which AI does a website use	Informational / Tool	Low	Homepage / Week 4 blog post
how to find what LLM a company uses	Informational	Low	Week 4 blog post

Notes: No established tools dominate these queries. AIHackr can rank #1 for all five within 60–90 days of publishing optimized content.

Cluster 3: Comparative — "OpenAI vs Anthropic Production Usage" (5 queries)

Keyword	Search Intent	Difficulty	Target Page
openai vs anthropic production usage	Research	Medium	Week 2 blog post
anthropic vs openai enterprise	Research	Medium	Week 7 blog post

Keyword	Search Intent	Difficulty	Target Page
openai vs anthropic enterprise saas	Research	Medium	Week 2 blog post
which llm do companies use in production	Research	Low–Medium	Week 2 blog post + `/leaderboard`
claude vs gpt enterprise deployment	Research	Medium	Week 7 blog post

Notes: These are higher-competition queries (OpenAI and Anthropic themselves rank) but AIHackr's unique angle — actual external detection data, not opinionated comparison — is a differentiated answer that can earn featured snippets.

Cluster 4: Informational — "How to Detect Which LLM" (5 queries)

Keyword	Search Intent	Difficulty	Target Page
how to detect which LLM a website uses	How-to	Low	Week 4 blog post
how to tell if a website uses ChatGPT	How-to	Low	Week 4 blog post
LLM fingerprinting	Technical / how-to	Low	Week 4 blog post
how to identify ai provider from network requests	Technical	Low	Week 4 blog post
detect openai api usage	Technical	Low	Week 4 blog post

Notes: Highly specific technical queries. Week 4's methodology post should dominate this cluster. Include code examples — code-heavy content ranks well for developer search intent.

Cluster 5: Leaderboard / Research — "YC Companies Using OpenAI" (5 queries)

Keyword	Search Intent	Difficulty	Target Page
yc companies using openai	Research	Low	`/leaderboard` (YC segment) + Week 5 blog post
yc ai companies llm provider	Research	Low	Week 5 blog post + `/leaderboard`
ai saas tech stack	Research	Medium	`/leaderboard` + Week 12 quarterly report
saas companies using anthropic	Research	Low	`/leaderboard` + `/provider/anthropic`
top saas ai provider breakdown	Research	Low	`/leaderboard` + Week 12 quarterly report

Notes: These research-intent queries are perfectly served by the leaderboard page. The `/leaderboard` URL should be the primary ranking target for all five — blog posts feed the leaderboard page via internal links.

SECTION 5: INTERNAL LINKING ARCHITECTURE

Link Graph (Text Representation)

```
HOMEPAGE (aihackr.com/)
├── → /stack (AI Stack Explorer index)
│       Text: "Explore AI stacks of 500+ companies"
├── → /leaderboard
│       Text: "View the AI Stack Leaderboard"
├── → /blog
│       Text: "AI stack research and data"
├── → /stack/[top 6 companies by traffic] (direct links in homepage hero)
│       Text: "[Company]'s AI stack →"
└── → Blog post Week 4 (methodology)
        Text: "How we detect AI providers"

/stack (index page — lists all company pages)
├── → /stack/[every company slug] (obvious — main purpose)
├── → /leaderboard
│       Text: "See how these companies rank on the AI Stack Leaderboard"
├── → Homepage scanner
│       Text: "Don't see your company? Run a scan →"
└── → Blog post Week 9 (AI stack checker comparison)
        Text: "How AIHackr compares to manual detection"

/stack/[company-slug] (individual company stack page)
├── → /leaderboard
│       Text: "See [Company] on the AI Stack Leaderboard"
├── → /provider/[detected-provider] (e.g., /provider/anthropic)
│       Text: "Other companies using Anthropic"
├── → /stack/[3-6 similar companies] (same provider, same category)
│       Text: "[Company]'s AI stack", "[Company]'s AI stack", etc.
├── → /stack (parent index)
│       Text: "← All company AI stacks" (breadcrumb)
├── → Blog post targeting same company (if exists)
│       Text: "Read our analysis of [Company]'s AI stack"
└── → Homepage scanner
        Text: "Run a scan on any domain"

/leaderboard (main leaderboard page)
├── → /stack/[every company in the table] (company name links in table)
├── → /provider/[each detected provider]
│       Text: "All companies using [Provider]"
├── → Blog post Week 2 (OpenAI vs Anthropic data)
│       Text: "See the full provider comparison →"
├── → Blog post Week 5 (YC companies)
│       Text: "YC AI company breakdown →"
├── → Blog post Week 11 (provider switching patterns)
│       Text: "Why companies switch LLM providers →"
├── → Blog post Week 12 (quarterly report)
│       Text: "Q3 2025 AI Stack Report →"
└── → Homepage scanner
        Text: "Scan a company not on this list"

/provider/[provider-slug] (e.g., /provider/anthropic)
```

```
├── → /leaderboard (filtered to that provider)
│       Text: "See all [Provider] companies on the leaderboard"
├── → /stack/[all companies using that provider]
│       Text: "[Company]'s AI stack", etc.
└── → Blog post Week 7 (Anthropic rise) or Week 2 (comparison)
        Text: "Read our analysis of [Provider]'s enterprise adoption"

BLOG POSTS
├── Each post → /leaderboard
│       Text: "See the live AI Stack Leaderboard" (bottom CTA on every post)
├── Each post → /stack/[companies mentioned in post]
│       Text: "[Company]'s AI stack" (inline contextual)
├── Each post → /stack (index)
│       Text: "Explore all company AI stacks" (sidebar or bottom of post)
├── Data posts (Weeks 2, 7, 11, 12) → each other
│       Text: referencing the specific data claim ("as we found in our Q2 analysis..
├── How-to posts (Weeks 4, 9) → homepage scanner
│       Text: "Try it yourself: scan any domain for free"
└── Company-specific posts (Weeks 1, 3, 6, 8, 10) → that company's /stack/ page
        Text: "[Company]'s detection data →" (prominent, near top of post)
```

Anchor Text Rules

Link destination	Preferred anchor text	Avoid
`/stack/[company]`	"[Company]'s AI stack", "[Company]'s LLM stack", "which AI [Company] uses"	"click here", "this page", "more info"
`/leaderboard`	"AI Stack Leaderboard", "weekly AI provider rankings", "see the leaderboard"	"here", "leaderboard" (alone)
`/stack` (index)	"AI stack explorer", "all company AI stacks", "browse AI stacks"	"page", "list"
`/provider/anthropic`	"all companies using Anthropic", "Anthropic's SaaS customers"	"Anthropic" (alone, too generic)
Homepage scanner	"scan any domain", "run a free scan", "try the AI detector"	"click here", "our tool"
Blog posts	Use the post's exact H1 or a natural excerpt	Truncated titles, "read more"

Page Authority Flow Priority

Highest authority pages (receive most inbound links, pass most link equity):

1. `/leaderboard` — the flagship, should earn the most external backlinks
2. `/stack/notion`, `/stack/intercom`, `/stack/hubspot`, `/stack/salesforce`, `/stack/cursor` — highest-traffic company pages
3. Blog post Week 2 (OpenAI vs Anthropic data) and Week 12 (quarterly report) — the two posts most likely to earn external links

These high-authority pages should link down to:

- Long-tail company stack pages (passes authority to pages that need it most)
- Provider pages (distributes across the provider taxonomy)

Lower-authority pages (individual company stack pages beyond top 20) should link up to:

- `/leaderboard`
- `/stack` index
- Thematically relevant blog posts

This creates a hub-and-spoke model where the leaderboard and top company pages are the hubs, and the long-tail company pages are the spokes.

Implementation Checklist for Monday Morning

Developer tasks (Day 1–2):

- [ ] Build `/stack/[slug]` template with all 10 data fields, evidence section, history timeline, and related company section
- [ ] Set up DB schema: `companies` table with all required fields
- [ ] Add JSON-LD schema (TechArticle + BreadcrumbList) to template
- [ ] Implement `sitemap.xml` auto-generation that includes all `/stack/` pages
- [ ] Set up `/stack` index page with search/filter UI
- [ ] Set up `/leaderboard` page with table, filters, and "This week's changes" panel
- [ ] Build OG image auto-generation for `/leaderboard/[YYYY-Www]/og.png`

Marketer tasks (Day 1–2):

- [ ] Populate the 50-company seed database with known data (manual research)
- [ ] Schedule Week 1 blog post (Notion) for Tuesday
- [ ] Create Twitter/X thread template in social scheduler
- [ ] Set up Monday 8am recurring calendar block for leaderboard review + publish
- [ ] Submit aihackr.com/sitemap.xml to Google Search Console
- [ ] Request indexing for all 50 `/stack/` pages via GSC URL Inspection

Day 7 (after first week live):

- [ ] Check GSC for impressions on `/stack/` pages — identify which queries are already surfacing
- [ ] Submit leaderboard URL to Product Hunt (Show PH: AI Stack Leaderboard)
- [ ] Post first HN submission with "Show HN" format

- [ ] Send cold pitch email to 3 newsletter authors with Week 2 data post as the hook

End of Part 3 — Content & Leaderboard Engine Playbook

For questions on execution, contact: strategy@aihackr.com